Chapter 7

- 7.1 Measures of predictive accuracy
- 7.2 Information criteria and cross-validation
  - Instead of 7.2, read:
    Vehtari, A., Gelman, A., Gabry, J. (2017). Practical Bayesian model evaluation using leave-one-out cross-validation and WAIC. Statistics and Computing. 27(5):1413–1432. arXiv preprint.
- 7.3 Model comparison based on predictive performance
- 7.4 Model comparison using Bayes factors
- 7.5 Continuous model expansion / sensitivity analysis
- 7.5 Example (may be skipped)

Model assessment, selection and inference after selection

- Extra material at https://avehtari.github.io/modelselection/
  - Videos, Slides, Notebooks, References

- CV-FAQ
  https://avehtari.github.io/modelselection/CV-FAQ.html

Predicting concrete quality

Predicting cancer recurrence
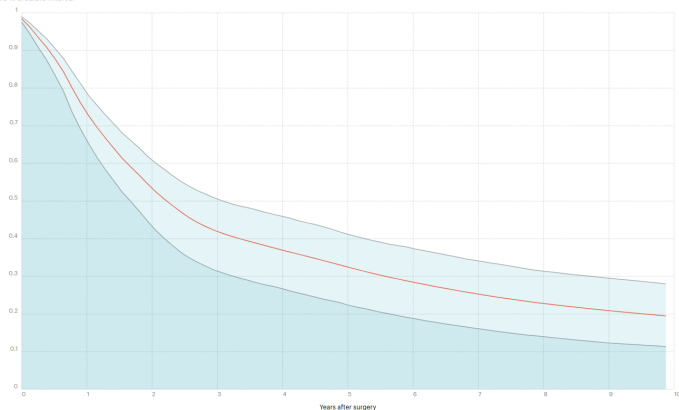
Predictive performance

- ▶ True predictive performance is found out by using it to make predictions and comparing predictions to true observations
  - ▶ external validation

Predictive performance

- ▶ True predictive performance is found out by using it to make predictions and comparing predictions to true observations
  - ▶ external validation
- ▶ Expected predictive performance
  - ▶ approximates the external validation

Predictive performance

- ▶ We need to choose the utility/cost function
- ▶ Application specific utility/cost functions are important
  - ▶ eg. money, life years, quality adjusted life years, etc.

Predictive performance

- ▶ We need to choose the utility/cost function
- ▶ Application specific utility/cost functions are important
  - ▶ eg. money, life years, quality adjusted life years, etc.
- ▶ If are interested overall in the goodness of the predictive distribution, or we don't know (yet) the application specific utility, then good information theoretically justified choice is log-score

$$\log p(y^{\text{rep}} \mid y, M),$$

- What is cross-validation
  - Leave-one-out cross-validation (elpd_loo, p_loo)
  - Uncertainty in LOO (SE)

- When is cross-validation applicable?
  - data generating mechanisms and prediction tasks
  - leave-many-out cross-validation

- Fast cross-validation
  - PSIS and diagnostics in loo package (Pareto k, n_eff, Monte Carlo SE)
  - K-fold cross-validation

- Related methods (WAIC, *IC, BF)

- Model comparison and selection (elpd_diff, se)

- Model averaging with Bayesian stacking

Stan and `loo` package

```
  Computed from 4000 by 20 log−likelihood matrix

          Estimate  SE
elpd_loo    −29.5  3.3
p_loo         2.7  1.0
──────
Monte Carlo SE of elpd_loo is 0.1.

Pareto k diagnostic values:
                          Count  Pct.    Min. n_eff
(−Inf, 0.5]   (good)       18    90.0%   899
 (0.5, 0.7]   (ok)          2    10.0%   459
   (0.7, 1]   (bad)         0     0.0%   <NA>
    (1, Inf)  (very bad)    0     0.0%   <NA>

All Pareto k estimates are ok (k < 0.7).
See help('pareto−k−diagnostic') for details.

Model comparison:
(negative 'elpd_diff' favors 1st model, positive favors 2nd)

elpd_diff        se
    −0.2        0.1
```
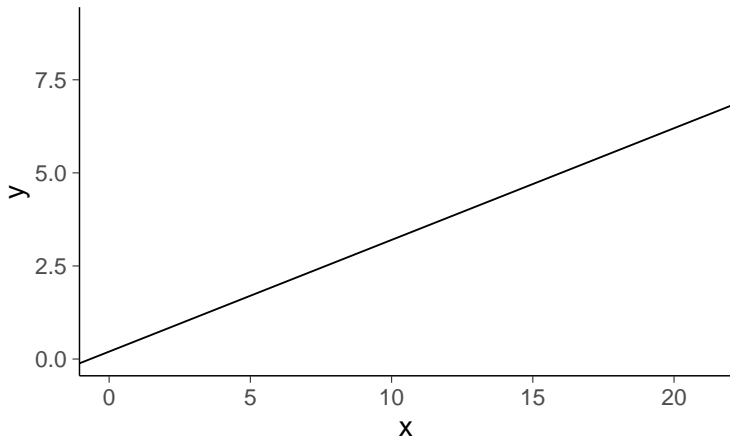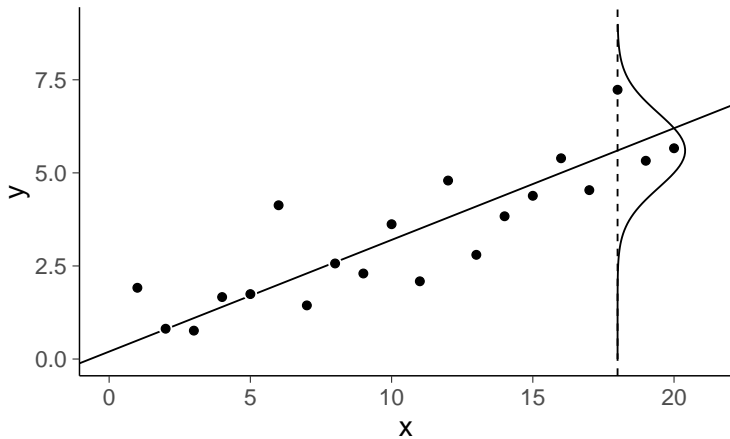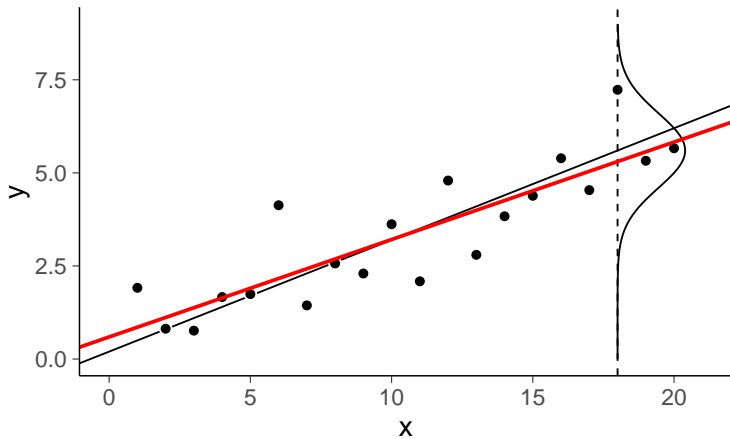
True mean y = a + bx

True mean and sigma
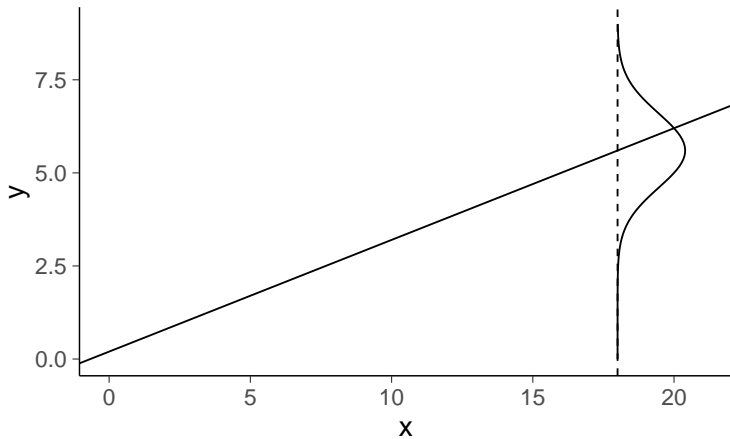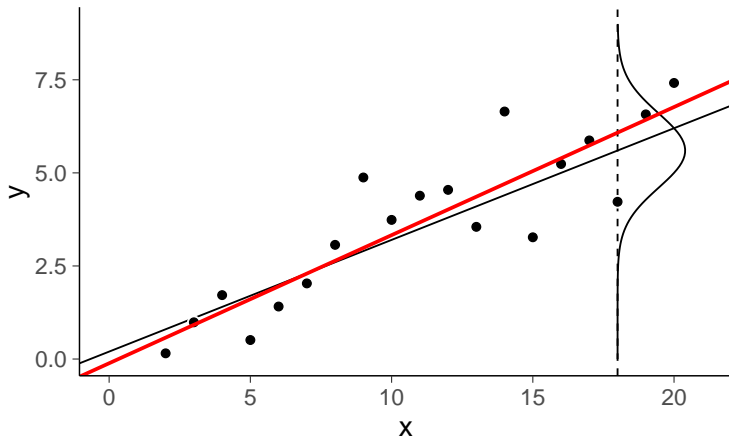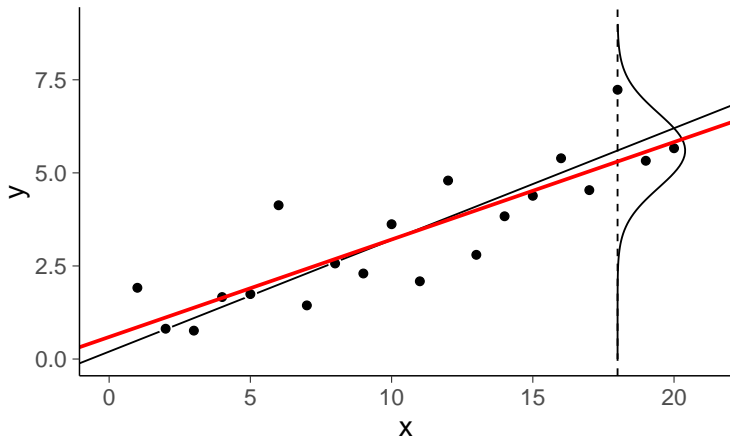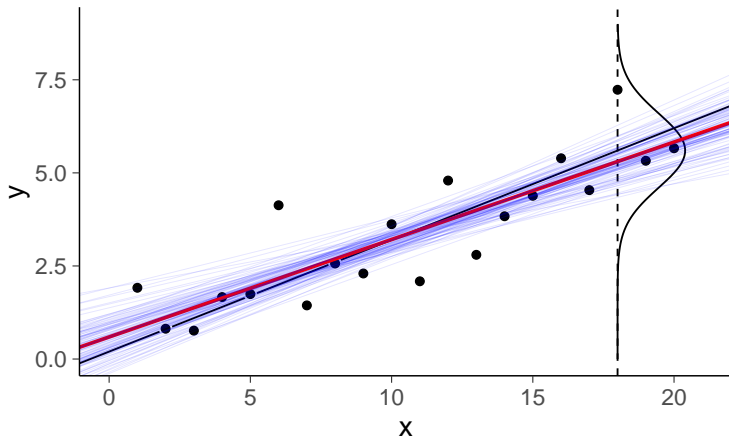
Data

Posterior mean

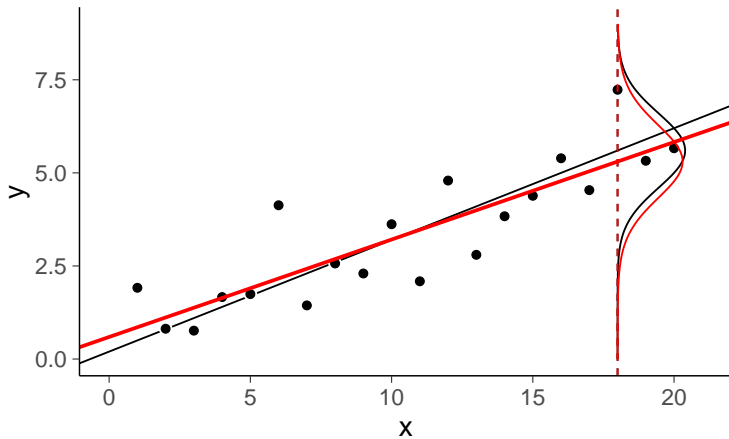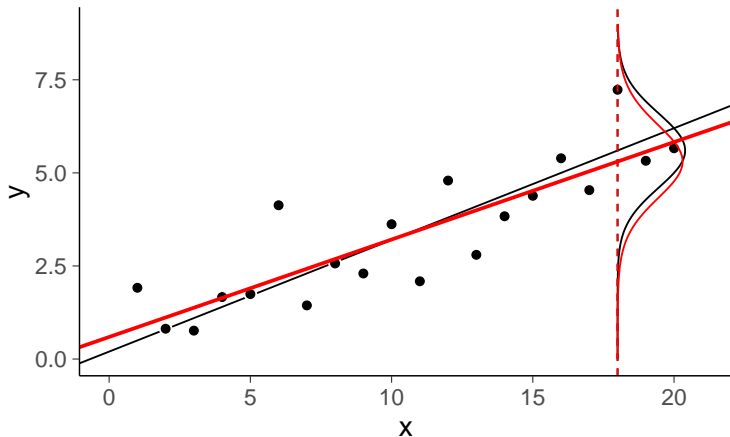Posterior mean, alternative data realisation

Posterior mean
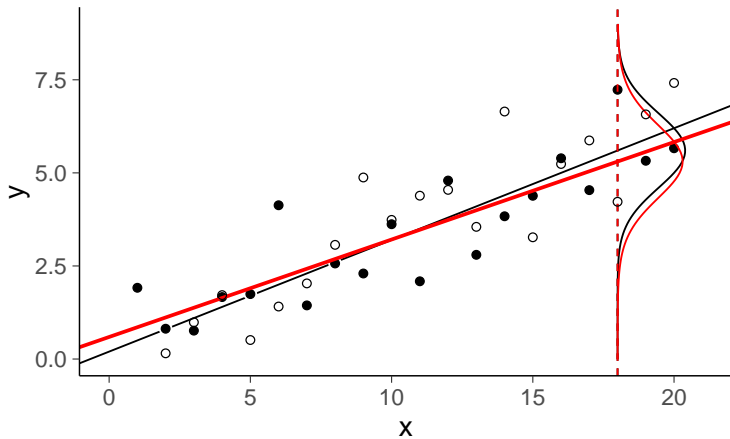
Posterior draws

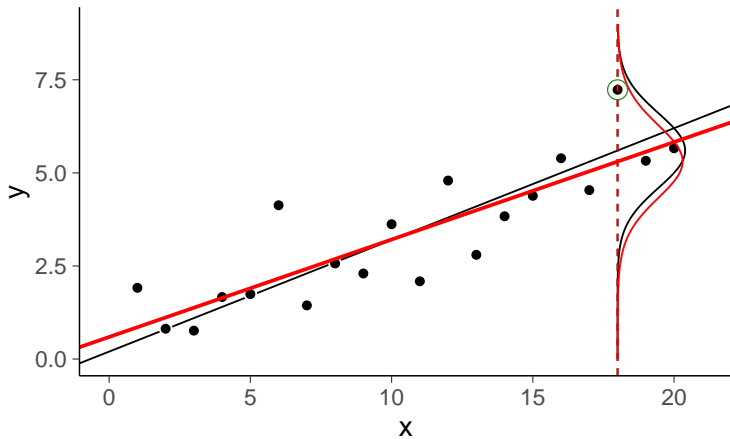# Posterior predictive distribution

Posterior predictive distribution

$$p(\tilde{y} \mid \tilde{x} = 18, x, y) = \int p(\tilde{y} \mid \tilde{x} = 18, \theta) p(\theta \mid x, y) d\theta$$
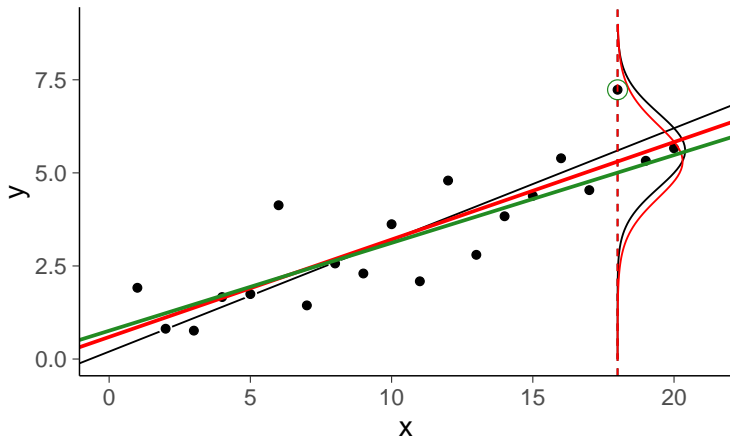
New data
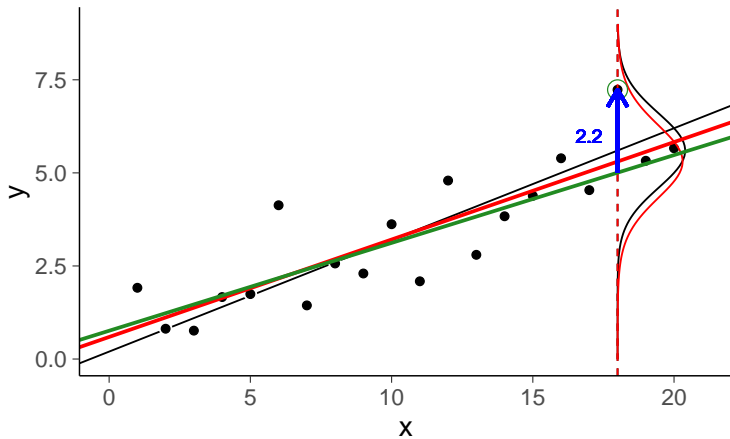
Posterior predictive distribution
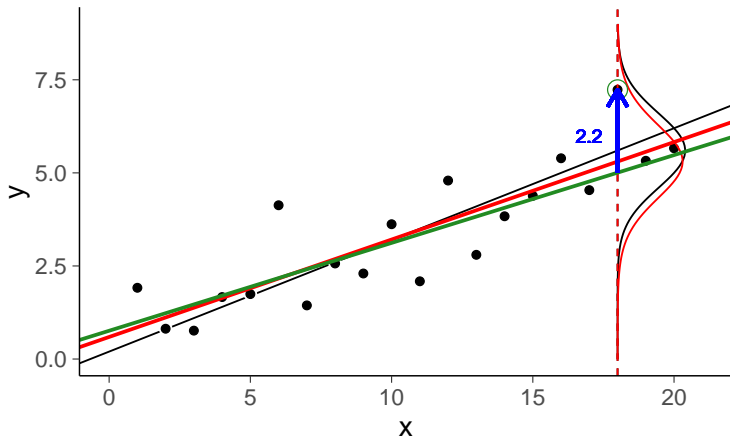
Leave−one−out mean

Leave−one−out residual

Leave−one−out residual

$y_{18} - E[p(\tilde{y} \mid \tilde{x} = 18, x_{-18}, y_{-18})]$

Leave−one−out residual
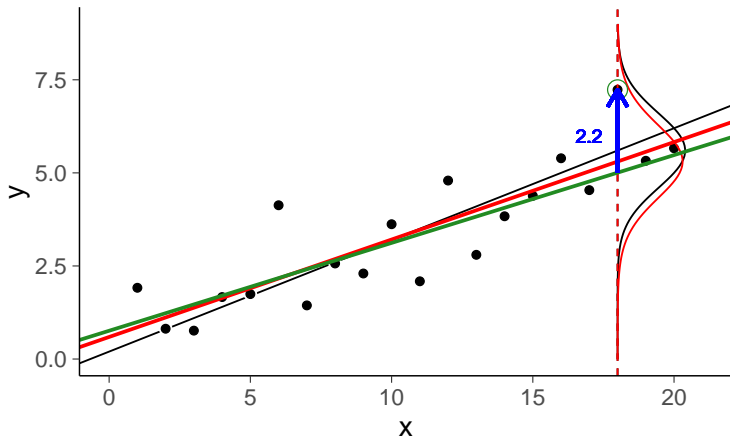
$y_{18} - E[p(\tilde{y} \mid \tilde{x} = 18, x_{-18}, y_{-18})]$

Can be use to compute, e.g., RMSE, R^2, 90% error

Leave−one−out residual
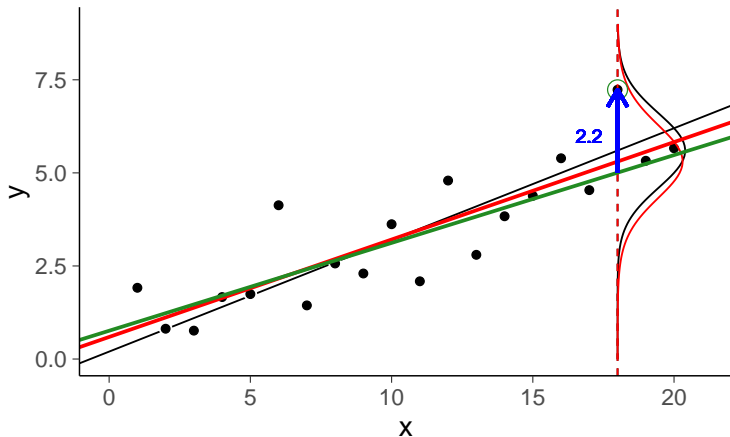
$y_{18} - E[p(\tilde{y} \mid \tilde{x} = 18, x_{-18}, y_{-18})]$

Can be use to compute, e.g., RMSE, R^2, 90% error

See LOO-R^2 at avehtari.github.io/bayes_R2/bayes_R2.html

Leave−one−out predictive distribution
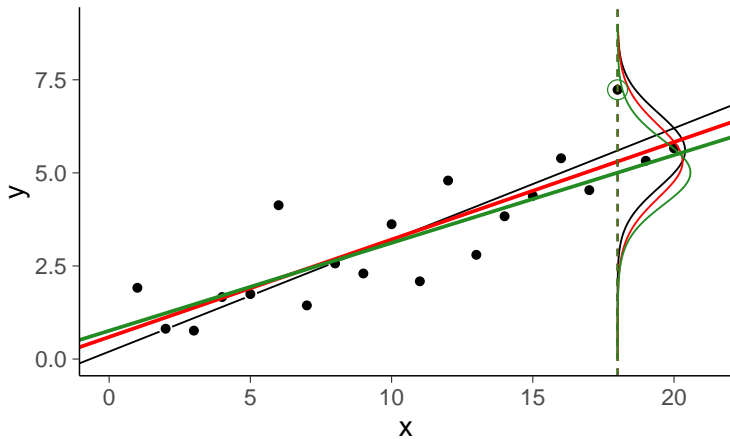
Leave−one−out predictive distribution
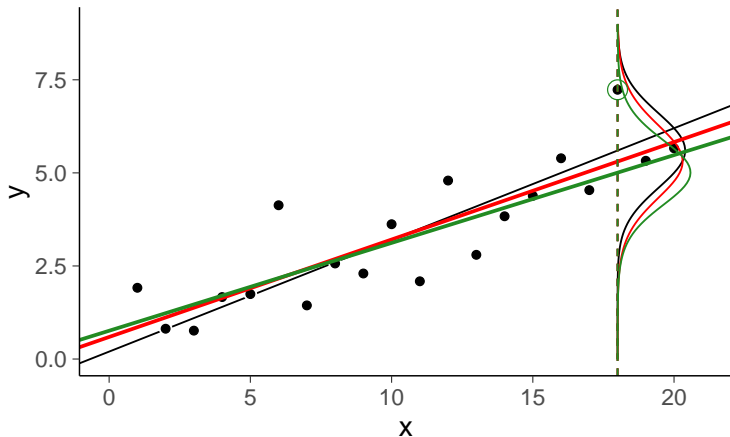
$$p(\tilde{y} \mid \tilde{x} = 18, x_{-18}, y_{-18}) = \int p(\tilde{y} \mid \tilde{x} = 18, \theta) p(\theta \mid x_{-18}, y_{-18}) d\theta$$

Posterior predictive density
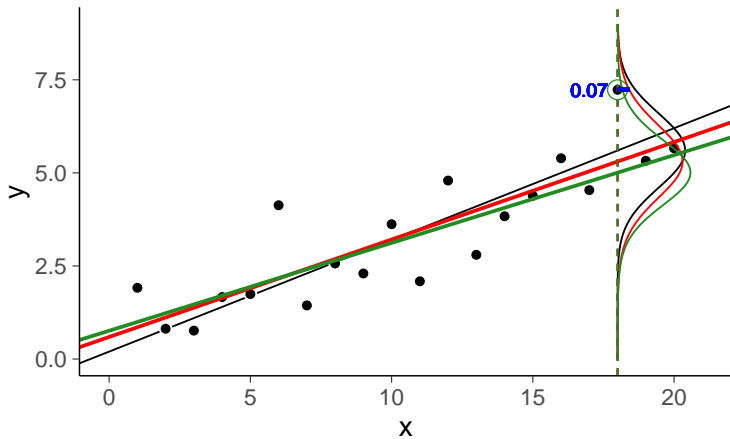
Posterior predictive density
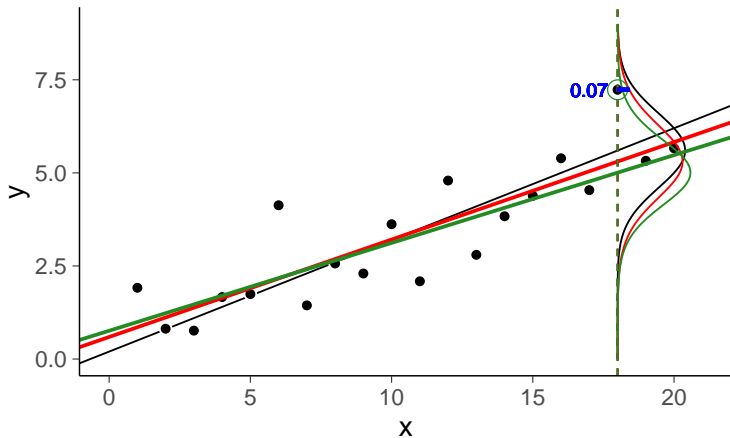
$$p(\tilde{y} = y_{18} \mid \tilde{x} = 18, x, y) \approx 0.07$$

Leave−one−out predictive density
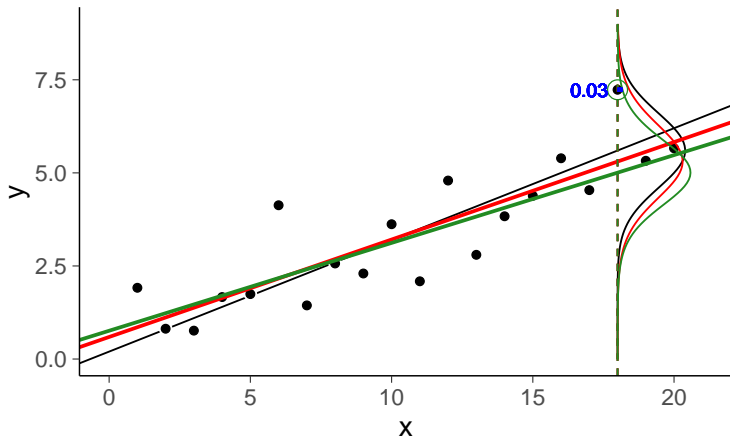
$p(\tilde{y} = y_{18} \mid \tilde{x} = 18, x, y) \approx 0.07$

$p(\tilde{y} = y_{18} \mid \tilde{x} = 18, x_{-18}, y_{-18}) \approx 0.03$

Leave−one−out predictive densities
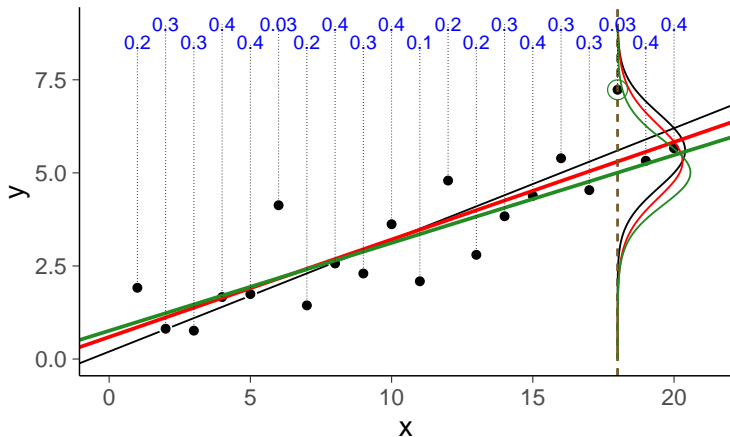
$p(y_i \mid x_i, x_{-i}, y_{-i}), \quad i = 1, \dots, 20$

Leave−one−out log predictive densities
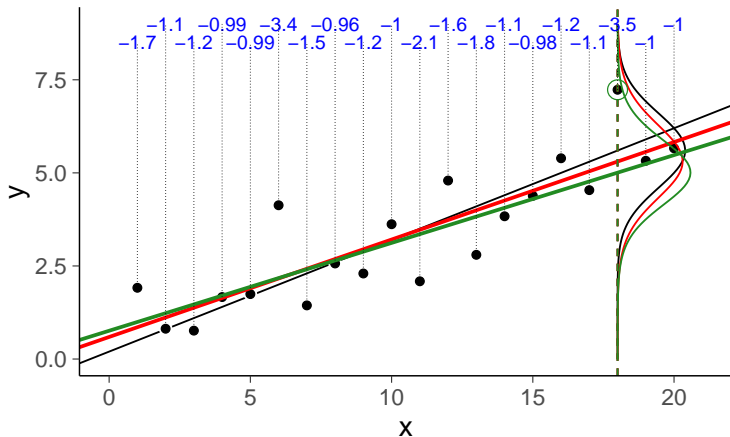
$\log p(y_i \mid x_i, x_{-i}, y_{-i}), \quad i = 1, \ldots, 20$

Leave−one−out log predictive densities
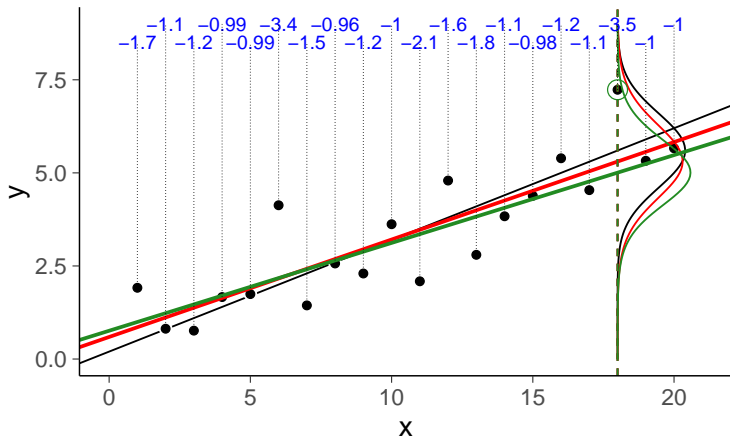
$$\sum_{i=1}^{20} \log p(y_i \mid x_i, x_{-i}, y_{-i}) \approx -29.5$$
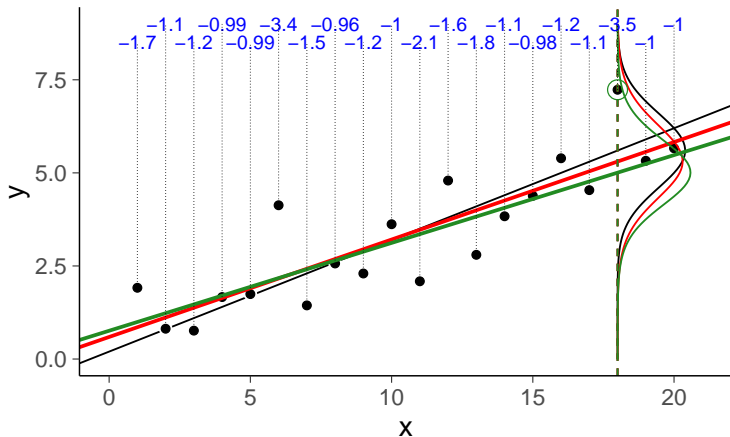
Leave−one−out log predictive densities

$$\text{elpd\_loo} = \sum_{i=1}^{20} \log p(y_i \mid x_i, x_{-i}, y_{-i}) \approx -29.5$$

Leave−one−out log predictive densities
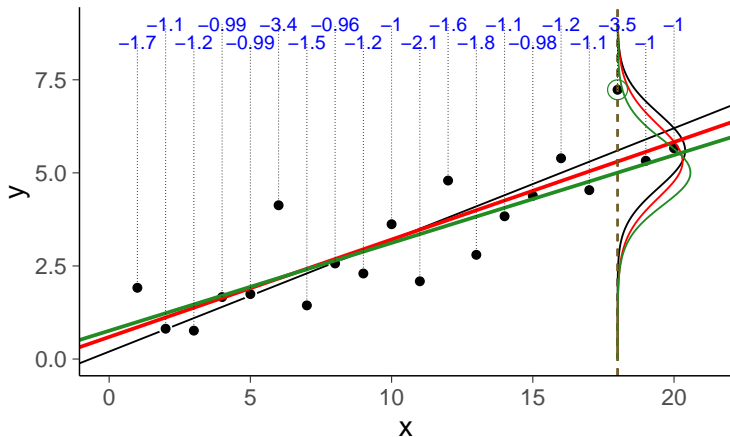
$$\text{elpd\_loo} = \sum_{i=1}^{20} \log p(y_i \mid x_i, x_{-i}, y_{-i}) \approx -29.5$$

an estimate of log posterior pred. density for new data

Leave−one−out log predictive densities

$$\text{elpd\_loo} = \sum_{i=1}^{20} \log p(y_i \mid x_i, x_{-i}, y_{-i}) \approx -29.5$$

$$\text{lpd} = \sum_{i=1}^{20} \log p(y_i \mid x_i, x, y) \approx -26.8$$

Leave−one−out log predictive densities
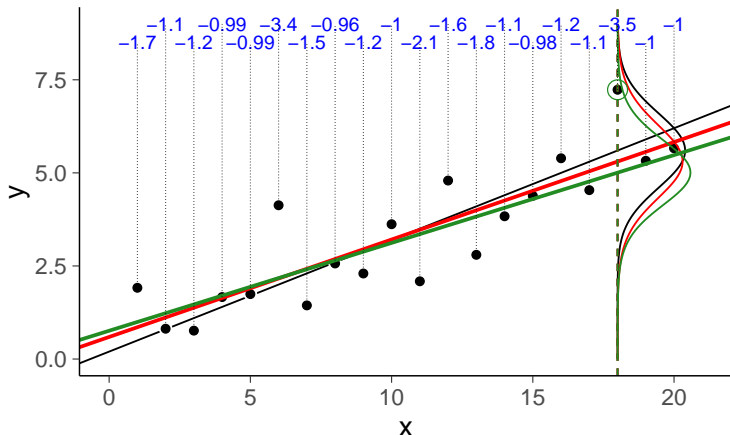
$\text{elpd\_loo} = \sum_{i=1}^{20} \log p(y_i \mid x_i, x_{-i}, y_{-i}) \approx -29.5$

$\text{lpd} = \sum_{i=1}^{20} \log p(y_i \mid x_i, x, y) \approx -26.8$

$\text{p\_loo} = \text{lpd} - \text{elpd\_loo} \approx 2.7$

Leave−one−out log predictive densities
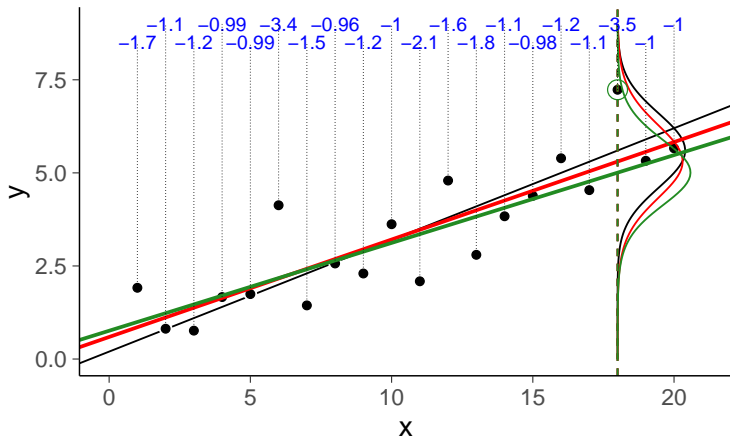
$$\text{elpd\_loo} = \sum_{i=1}^{20} \log p(y_i \mid x_i, x_{-i}, y_{-i}) \approx -29.5$$

$$\text{SE} = \text{sd}(\log p(y_i \mid x_i, x_{-i}, y_{-i})) \cdot \sqrt{20} \approx 3.3$$

Leave−one−out log predictive densities
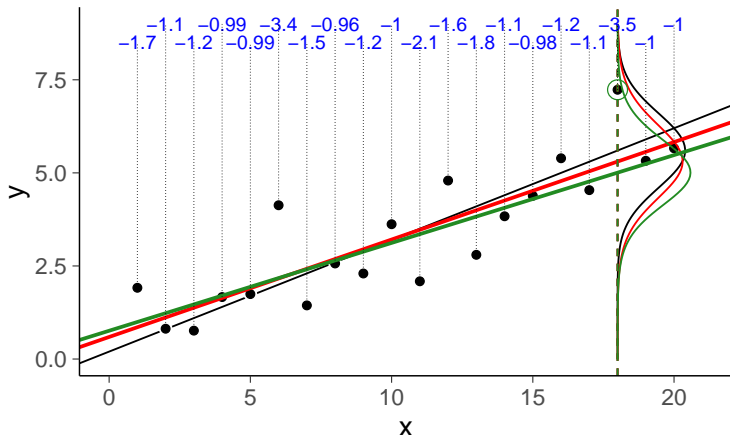
$$\text{elpd\_loo} = \sum_{i=1}^{20} \log p(y_i \mid x_i, x_{-i}, y_{-i}) \approx -29.5$$
$$\text{SE} = \text{sd}(\log p(y_i \mid x_i, x_{-i}, y_{-i})) \cdot \sqrt{20} \approx 3.3$$

see Vehtari, Gelman & Gabry (2017a) and Vehtari & Ojanen (2012) for more

Fixed / designed x
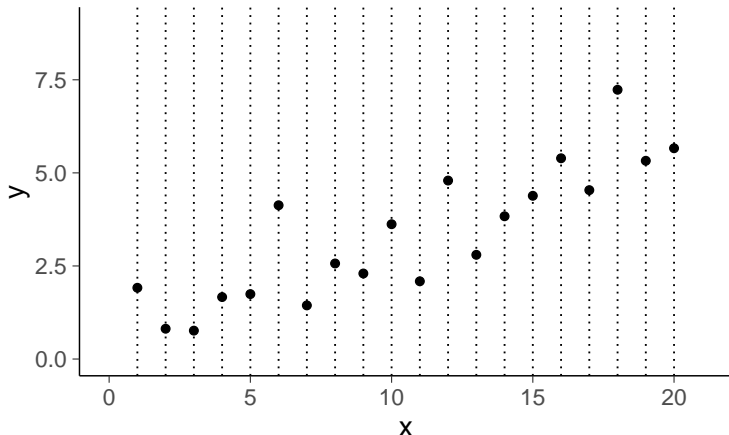
LOO is ok for fixed / designed *x*. SE is uncertainty about *y* | *x*.

see Vehtari & Ojanen (2012) and CV-FAQ
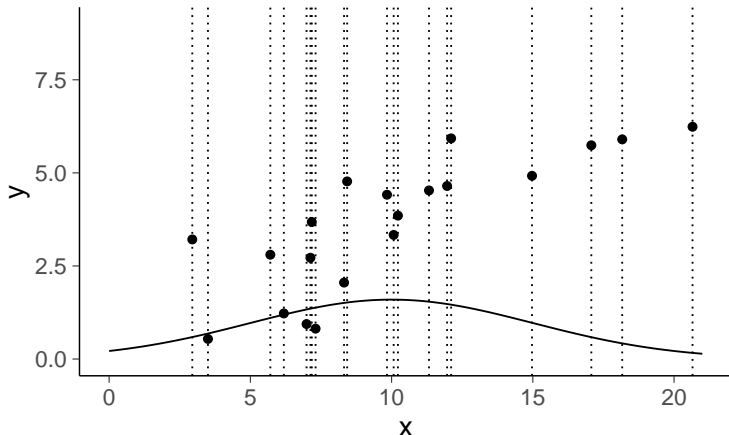
Distribution for x

LOO is ok for random *x*. SE is uncertainty about *y* | *x* and *x*.

see Vehtari & Ojanen (2012) and CV-FAQ

Distribution for x
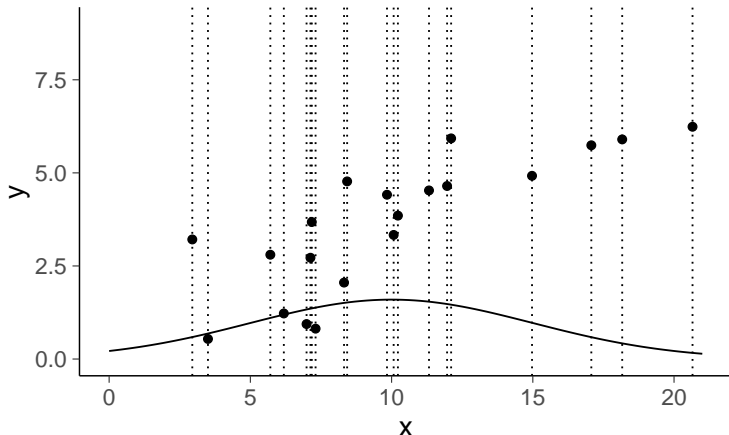
LOO is ok for random *x*. SE is uncertainty about $y \mid x$ and *x*.

Covariate shift can be handled with importance weighting or modelling

see Vehtari & Ojanen (2012) and CV-FAQ

loo package

Computed from 4000 by 20 log−likelihood matrix

```
         Estimate  SE
elpd_loo   −29.5   3.3
p_loo        2.7   1.0
```

Monte Carlo SE of elpd_loo is 0.1.

Pareto k diagnostic values:

		Count	Pct.	Min. n_eff
(−Inf, 0.5]	(good)	18	90.0%	899
(0.5, 0.7]	(ok)	2	10.0%	459
(0.7, 1]	(bad)	0	0.0%	<NA>
(1, Inf)	(very bad)	0	0.0%	<NA>

All Pareto k estimates are ok (k < 0.7).
See help('pareto−k−diagnostic') for details.

Interpolation vs extrapolation

Interpolation vs extrapolation

Nonlinear model fit

Interpolation vs extrapolation

Nonlinear model fit + new data

Interpolation vs extrapolation

Nonlinear model fit + new data

Extrapolation is more difficult

Cross-validation for time series?

Can LOO or other cross-validation be used with time series?

Cross-validation for time series

Leave-one-out cross-validation is ok for assessing conditional model

Cross-validation for time series

Leave-future-out cross-validation is better for predicting future

Cross-validation for time series

m-step-ahead cross-validation is better for predicting further future

Cross-validation for time series

m-step-ahead leave-a-block-out cross-validation

Cross-validation for hierarchical data

Rats data

Can LOO or other cross-validation be used with hierarchical data?

Cross-validation for hierarchical data

Leave−one−out?

Yes!

Cross-validation for hierarchical data

1−step−ahead?

Yes!

Cross-validation for hierarchical data

Leave−one−time−point−out?

Yes!

Cross-validation for hierarchical data

Leave−one−rat−out?

Yes!

Cross-validation for hierarchical data

Predict given initial weight?

Yes!

Summary of data generating mechanisms and prediction tasks

- You have to make some assumptions on data generating mechanism

- Use the knowledge of the prediction task if available

- Cross-validation can be used to analyse different parts, even if there is no clear prediction task

see Vehtari & Ojanen (2012) and CV-FAQ

Fast cross-validation

- Pareto smoothed importance sampling LOO (PSIS-LOO)
- K-fold cross-validation

see Vehtari, Gelman & Gabry (2017a) and mc-stan.org/loo/

Importance sampling leave-one-out cross-validation

- We want to compute
  $p(y_i \mid x_{-i}, y_{-i}) = \int p(y_i \mid x_i, \theta) p(\theta \mid x_{-i}, y_{-i}) d\theta$

Importance sampling leave-one-out cross-validation

- We want to compute
  $p(y_i \mid x_{-i}, y_{-i}) = \int p(y_i \mid x_i, \theta) p(\theta \mid x_{-i}, y_{-i}) d\theta$
- Proposal distribution is full posterior $\theta^{(s)} \sim p(\theta \mid x, y)$
- Target distribution is LOO-posterior $p(\theta \mid x_{-i}, y_{-i})$

Importance sampling leave-one-out cross-validation

- We want to compute
  $p(y_i \mid x_{-i}, y_{-i}) = \int p(y_i \mid x_i, \theta) p(\theta \mid x_{-i}, y_{-i}) d\theta$
- Proposal distribution is full posterior $\theta^{(s)} \sim p(\theta \mid x, y)$
- Target distribution is LOO-posterior $p(\theta \mid x_{-i}, y_{-i})$
- Importance ratio

$$w_i^{(s)} = \frac{p(\theta^{(s)} \mid x_{-i}, y_{-i})}{p(\theta^{(s)} \mid x, y)} \propto \frac{1}{p(y_i \mid \theta^{(s)})}$$

Importance sampling leave-one-out cross-validation

- We want to compute
  $p(y_i \mid x_{-i}, y_{-i}) = \int p(y_i \mid x_i, \theta) p(\theta \mid x_{-i}, y_{-i}) d\theta$

- Proposal distribution is full posterior $\theta^{(s)} \sim p(\theta \mid x, y)$

- Target distribution is LOO-posterior $p(\theta \mid x_{-i}, y_{-i})$

- Importance ratio

$$w_i^{(s)} = \frac{p(\theta^{(s)} \mid x_{-i}, y_{-i})}{p(\theta^{(s)} \mid x, y)} \propto \frac{1}{p(y_i \mid \theta^{(s)})}$$

$$\tilde{w}_i^{(s)} = \frac{w_i^{(s)}}{\sum_{s'=1}^{S} w_i^{(s')}}$$

Data

Posterior draws

$$\theta^{(s)} \sim p(\theta \mid x, y)$$

Posterior predictive distribution

$$\theta^{(s)} \sim p(\theta \mid x, y), \quad p(\tilde{y} \mid \tilde{x}, x, y) \approx \frac{1}{S} \sum_{s=1}^{S} p(\tilde{y} \mid \tilde{x}, \theta^{(s)})$$

Posterior predictive distribution

$$\theta^{(s)} \sim p(\theta \mid x, y), \quad p(\tilde{y} \mid \tilde{x}, x, y) \approx \frac{1}{S} \sum_{s=1}^{S} p(\tilde{y} \mid \tilde{x}, \theta^{(s)})$$

PSIS−LOO weighted draws

$$\theta^{(s)} \sim p(\theta \mid x, y), \quad w_i^{(s)} = p(\theta^{(s)} \mid x_{-i}, y_{-i}) / p(\theta^{(s)} \mid x, y)$$

PSIS−LOO weighted predictive distribution

$$\theta^{(s)} \sim p(\theta \mid x, y), \quad w_i^{(s)} = p(\theta^{(s)} \mid x_{-i}, y_{-i}) / p(\theta^{(s)} \mid x, y)$$

$$p(y_i \mid x_i, x_{-i}, y_{-i}) \approx \sum_{s=1}^{S} [\tilde{w}_i^{(s)} p(y_i \mid x_i, \theta^{(s)})]$$

Pareto smoothed importance sampling LOO

- We want to compute
  $p(y_i \mid x_{-i}, y_{-i}) = \int p(y_i \mid x_i, \theta) p(\theta \mid x_{-i}, y_{-i}) d\theta$

- Proposal distribution is full posterior $\theta^{(s)} \sim p(\theta \mid x, y)$

- Target distribution is LOO-posterior $p(\theta \mid x_{-i}, y_{-i})$

- Importance ratio

$$w_i^{(s)} = \frac{p(\theta^{(s)} \mid x_{-i}, y_{-i})}{p(\theta^{(s)} \mid x, y)} \propto \frac{1}{p(y_i \mid \theta^{(s)})}$$

$$\tilde{w}_i^{(s)} = \frac{w_i^{(s)}}{\sum_{s'=1}^{S} w_i^{(s')}}$$

- The variability of importance weights matter
    - Pareto-*k* diagnostic
    - Pareto smoothed importance sampling LOO
      (PSIS-LOO)

400 importance weights for leave−18th−out

4000 importance weights for leave−18th−out

Equal weights

w_i

4000 importance weights for leave−18th−out

Equal weights

w_i

$\text{ESS} \approx 1/\sum_{s=1}^{S}(\tilde{w}^{(s)})^2 \approx 459$

see Vehtari, Gelman & Gabry (2017b)

4000 importance weights for leave−18th−out

ESS $\approx 1/\sum_{s=1}^{S}(\tilde{w}^{(s)})^2 \approx 459$

Pareto $\hat{k} \approx 0.52$

- Pareto $\hat{k}$ estimates the tail shape which determines the convergence rate of PSIS. Less than 0.7 is ok.

see Vehtari, Gelman & Gabry (2017b)

PSIS–LOO diagnostics

PSIS–LOO diagnostics

Pareto k diagnostic values:

		Count	Pct.	Min. n_eff
(−Inf, 0.5]	(good)	18	90.0%	899
(0.5, 0.7]	(ok)	2	10.0%	459
(0.7, 1]	(bad)	0	0.0%	<NA>
(1, Inf)	(very bad)	0	0.0%	<NA>

PSIS−LOO diagnostics

```
          Pareto k diagnostic values:
                              Count  Pct.    Min. n_eff
(−Inf , 0.5]   (good)          18    90.0%    899
 (0.5 , 0.7]   (ok)             2    10.0%    459
   (0.7 , 1]   (bad)            0     0.0%    <NA>
   (1 , Inf)   (very bad)       0     0.0%    <NA>
```

loo package

```
          Estimate  SE
elpd_loo    −29.5  3.3
p_loo         2.7  1.0
```

_____

Monte Carlo SE of elpd_loo is 0.1.

Pareto k diagnostic values:

		Count	Pct.	Min. n_eff
(−Inf, 0.5]	(good)	18	90.0%	899
(0.5, 0.7]	(ok)	2	10.0%	459
(0.7, 1]	(bad)	0	0.0%	<NA>
(1, Inf)	(very bad)	0	0.0%	<NA>

All Pareto k estimates are ok (k < 0.7).
See help('pareto−k−diagnostic') for details.

see more in Vehtari, Gelman & Gabry (2017b)

Pareto smoothed importance sampling (PSIS)

- Replace the largest weights with ordered statistics of the fitted Pareto distribution
    - equivalent to using model to filter the noise out of the weights

See more in Vehtari, Simpson, Gelman, Yao & Gabry (2021)

Pareto smoothed importance sampling (PSIS)

- Replace the largest weights with ordered statistics of the fitted Pareto distribution
    - equivalent to using model to filter the noise out of the weights

- Reduced variability compared to the plain IS

- Reduced bias compared to the truncated IS

See more in Vehtari, Simpson, Gelman, Yao & Gabry (2021)

Pareto smoothed importance sampling (PSIS)

- Replace the largest weights with ordered statistics of the fitted Pareto distribution
    - equivalent to using model to filter the noise out of the weights

- Reduced variability compared to the plain IS

- Reduced bias compared to the truncated IS

- Asymptotically consistent under some mild conditions

See more in Vehtari, Simpson, Gelman, Yao & Gabry (2021)

Stan code

$$\log(w_i^{(s)}) = \log(1/p(y_i \mid x_i, \theta^{(s)})) = -\texttt{log\_lik[i]}$$

Stan code

$$\log(w_i^{(s)}) = \log(1/p(y_i \mid x_i, \theta^{(s)})) = -\texttt{log\_lik[i]}$$

```
...
model {
  alpha ~ normal(pmualpha, psalpha);
  beta ~ normal(pmubeta, psbeta);
  y ~ normal(mu, sigma);
}
generated quantities {
  vector[N] log_lik;
  for (i in 1:N)
    log_lik[i] = normal_lpdf(y[i] | mu[i], sigma);
}
```

Stan code

$$\log(w_i^{(s)}) = \log(1/p(y_i \mid x_i, \theta^{(s)})) = -\texttt{log\_lik[i]}$$

```
...
model {
  alpha ~ normal(pmualpha, psalpha);
  beta ~ normal(pmubeta, psbeta);
  y ~ normal(mu, sigma);
}
generated quantities {
  vector[N] log_lik;
  for (i in 1:N)
    log_lik[i] = normal_lpdf(y[i] | mu[i], sigma);
}
```

- RStanARM and BRMS compute `log_lik` by default

Pareto smoothed importance sampling LOO

- PSIS-LOO for hierarchical models
    - leave-one-group out is challenging for PSIS-LOO
    - see Merkel, Furr and Rabe-Hesketh (2018) for an approach using quadrature integration
    - Stan demo of the challenges and integrated LOO at https://avehtari.github.io/modelselection/roaches.html

Pareto smoothed importance sampling LOO

- PSIS-LOO for hierarchical models
  - leave-one-group out is challenging for PSIS-LOO
  - see Merkel, Furr and Rabe-Hesketh (2018) for an approach using quadrature integration
  - Stan demo of the challenges and integrated LOO at https://avehtari.github.io/modelselection/roaches.html
- PSIS-LOO for non-factorizable models
  - mc-stan.org/loo/articles/loo2-non-factorizable.html

Pareto smoothed importance sampling LOO

- PSIS-LOO for hierarchical models
  - leave-one-group out is challenging for PSIS-LOO
  - see Merkel, Furr and Rabe-Hesketh (2018) for an approach using quadrature integration
  - Stan demo of the challenges and integrated LOO at https://avehtari.github.io/modelselection/roaches.html
- PSIS-LOO for non-factorizable models
  - mc-stan.org/loo/articles/loo2-non-factorizable.html
- PSIS-LOO for time series
  - Approximate leave-future-out cross-validation (LFO-CV)
    mc-stan.org/loo/articles/loo2-lfo.html

Data

AR−4 prediction with 95% interval

LFO-CV

PSIS−1−step−ahead with refits

K-fold cross-validation

- K-fold cross-validation can approximate LOO
  - the same use cases as with LOO
- K-fold cross-validation can be used for hierarchical models
  - good for leave-one-group-out
- K-fold cross-validation can be used for time series
  - with leave-block-out

Balance k−fold approximation of LOO

Balance k−fold approximation of LOO

Random k−fold approximation of LOO

Random kfold approximation of LOO

Leave−one−rat−out
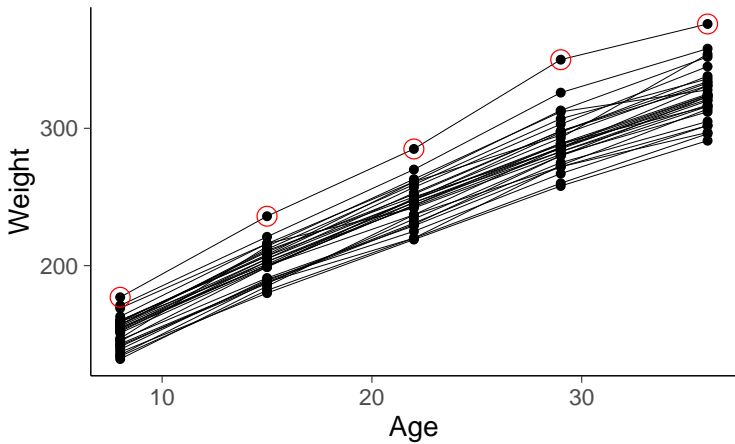
Leave−one−rat−out

kfold_split_random()
kfold_split_balanced()
kfold_split_stratified()

WAIC vs PSIS-LOO

see Vehtari, Gelman & Gabry (2017a)

WAIC vs PSIS-LOO

- WAIC has same assumptions as LOO

see Vehtari, Gelman & Gabry (2017a)

WAIC vs PSIS-LOO

- WAIC has same assumptions as LOO
- PSIS-LOO is more accurate

see Vehtari, Gelman & Gabry (2017a)

WAIC vs PSIS-LOO

- WAIC has same assumptions as LOO

- PSIS-LOO is more accurate

- PSIS-LOO has much better diagnostics

see Vehtari, Gelman & Gabry (2017a)

WAIC vs PSIS-LOO

- WAIC has same assumptions as LOO

- PSIS-LOO is more accurate

- PSIS-LOO has much better diagnostics

- LOO makes the prediction assumption more clear, which helps if K-fold-CV is needed instead

see Vehtari, Gelman & Gabry (2017a)

WAIC vs PSIS-LOO

- WAIC has same assumptions as LOO

- PSIS-LOO is more accurate

- PSIS-LOO has much better diagnostics

- LOO makes the prediction assumption more clear, which helps if K-fold-CV is needed instead

- Multiplying by -2 doesn't give any benefit (Watanabe didn't multiply by -2)

see Vehtari, Gelman & Gabry (2017a)

*IC

- AIC uses maximum likelihood estimate for prediction
- DIC uses posterior mean for prediction
- BIC is an approximation for marginal likelihood
- TIC, NIC, RIC, PIC, BPIC, QIC, AICc, ...

Marginal likelihood / Bayes factor

- Like leave-future-out 1-step-ahead cross-validation but starting with 0 observations

Marginal likelihood / Bayes factor

- Like leave-future-out 1-step-ahead cross-validation but starting with 0 observations

Marginal likelihood / Bayes factor

- Like leave-future-out 1-step-ahead cross-validation but starting with 0 observations
  - which makes it very sensitive to prior

Marginal likelihood / Bayes factor

- Like leave-future-out 1-step-ahead cross-validation but starting with 0 observations
  - which makes it very sensitive to prior and
  - unstable in case of misspecified models

Marginal likelihood / Bayes factor

- Like leave-future-out 1-step-ahead cross-validation but starting with 0 observations
  - which makes it very sensitive to prior and
  - unstable in case of misspecified models also asymptotically
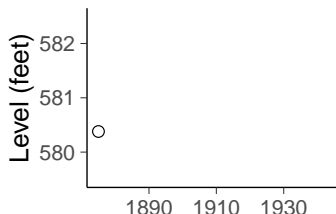
Cross-validation for model assessment

- CV is good for model assessment when application specific utility/cost functions are used
  - e.g. 90% absolute error

Cross-validation for model assessment

- CV is good for model assessment when application specific utility/cost functions are used
  - e.g. 90% absolute error

- Also useful in model checking in similar way as posterior predictive checking (PPC)
  - model misspecification diagnostics (e.g. Pareto-k and p_loo)
  - checking calibration of leave-one-out predictive posteriors (ppc_loo_pit in bayesplot)

  see demos avehtari.github.io/modelselection/

Radon example

PSIS-LOO diagnostics

see Vehtari, Gelman & Gabry (2017a)

Sometimes cross-validation is not needed

Sometimes cross-validation is not needed

- Posterior predictive checking is often sufficient

Model for weight

Model for log(weight)

Predicting the yields of mesquite bushes.

Gelman, Hill & Vehtari (2020): Regression and Other Stories, Chapter 11.

Sometimes cross-validation is not needed

- Posterior predictive checking is often sufficient

Model for weight

Model for log(weight)

Predicting the yields of mesquite bushes.

Gelman, Hill & Vehtari (2020): Regression and Other Stories, Chapter 11.

- BDA3, Chapter 6
- Gabry, Simpson, Vehtari, Betancourt, Gelman (2019). Visualization in Bayesian workflow. JRSS A, https://doi.org/10.1111/rssa.12378
- mc-stan.org/bayesplot/articles/graphical-ppcs.html

Sometimes cross-validation is not needed

- With good priors that keep the prior on predictive space consistent, there is no need to do model selection to avoid overfitting

Arsenic well example – Model comparison

- Logistic regression for predicting probability of switching well with high arsenic level in rural Bangladesh

    - Model 1:
      log(arsenic) + distance

    - Model 2:
      log(arsenic) + distance + education level

Arsenic well example – Model comparison

Model 1 vs Model 2

Model 1: $\widehat{\mathrm{elpd}}_{\mathrm{LOO}}\left(\mathrm{M}_a \mid y^{\mathrm{obs}}\right) \approx$ -1952, SE=16
Model 2: $\widehat{\mathrm{elpd}}_{\mathrm{LOO}}\left(\mathrm{M}_b \mid y^{\mathrm{obs}}\right) \approx$ -1938, SE=17

Arsenic well example – Model comparison

Model 1 vs Model 2

Difference: $\widehat{\text{elpd}}_{\text{LOO}}\left(\text{M}_a, \text{M}_b \mid y^{\text{obs}}\right) \approx$ -14.4, SE = 6.1

Arsenic well example – Model comparison

Model 1 vs Model 2

Difference: $\widehat{\mathrm{elpd}}_{\mathrm{LOO}}\left(\mathrm{M}_a, \mathrm{M}_b \mid y^{\mathrm{obs}}\right) \approx$ -14.4, SE = 6.1

Arsenic well example – Model comparison

Model 1 vs Model 2

Difference: $\widehat{\mathrm{elpd}}_{\mathrm{LOO}}\left(\mathrm{M}_a, \mathrm{M}_b \mid y^{\mathrm{obs}}\right) \approx$ -14.4, SE = 6.1

Arsenic well example – Model comparison

Difference: $\widehat{\text{elpd}}_{\text{LOO}}\left(M_a, M_b \mid y^{\text{obs}}\right) \approx$ -14.4, SE = 6.1

Arsenic well example – Model comparison


```
> loo_compare(model1, model2)
        elpd_diff  se_diff
model2    0.0        0.0
model1  -14.4        6.1
```

see Vehtari, Gelman & Gabry (2017a)

LOO difference uncertainty estimate

1. The models make very similar predictions

2. The models are misspecified with outliers in the data

3. The number of observations is small

LOO difference uncertainty estimate

1. The models make very similar predictions
   - if $|\widehat{\mathrm{elpd}}_{\mathrm{LOO}}(\mathrm{M}_a, \mathrm{M}_b \mid y^{\mathrm{obs}})| < 4$, SE is not reliable, but the difference is small anyway
   - selecting a "wrong" model has small cost
   - in nested case the skewness favors the simpler model

2. The models are misspecified with outliers in the data

3. The number of observations is small

LOO difference uncertainty estimate

1. The models make very similar predictions
   - if $|\widehat{\mathrm{elpd}}_{\mathrm{LOO}}(\mathrm{M}_a, \mathrm{M}_b \mid y^{\mathrm{obs}})| < 4$, SE is not reliable, but the difference is small anyway
   - selecting a "wrong" model has small cost
   - in nested case the skewness favors the simpler model

2. The models are misspecified with outliers in the data
   - in nested case the bias favors the simpler model
   - model checking and model extension to avoid misspecified models (Bayesian workflow)

3. The number of observations is small

LOO difference uncertainty estimate

1. The models make very similar predictions
   - if $|\widehat{\mathrm{elpd}}_{\mathrm{LOO}}(\mathrm{M}_a, \mathrm{M}_b \mid y^{\mathrm{obs}})| < 4$, SE is not reliable, but the difference is small anyway
   - selecting a "wrong" model has small cost
   - in nested case the skewness favors the simpler model

2. The models are misspecified with outliers in the data
   - in nested case the bias favors the simpler model
   - model checking and model extension to avoid misspecified models (Bayesian workflow)

3. The number of observations is small
   - in nested case the skewness favors the simpler model
   - any inference with small n is difficult
   - if $|\widehat{\mathrm{elpd}}_{\mathrm{LOO}}(\mathrm{M}_a, \mathrm{M}_b \mid y^{\mathrm{obs}})| > 4$, model is well specified, and $n > 100$ then the normal approximation is good

Sometimes cross-validation is not needed

- In nested case, often easier and more accurate to analyse posterior distribution of more complex model directly
  avehtari.github.io/modelselection/betablockers.html

Sometimes predictive model comparison can be useful
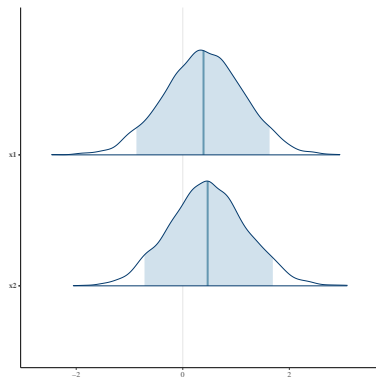
Marginal posterior intervals

Sometimes predictive model comparison can be useful
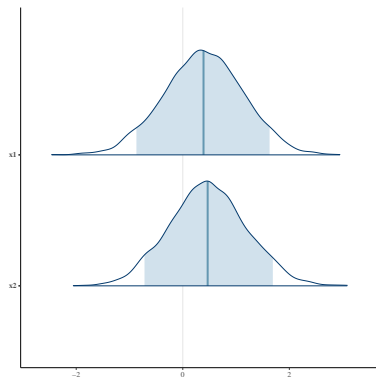
Marginal posterior intervals   Joint posterior density
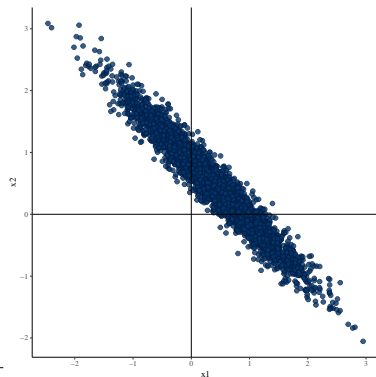
rstanarm + bayesplot

Sometimes predictive model comparison can be useful
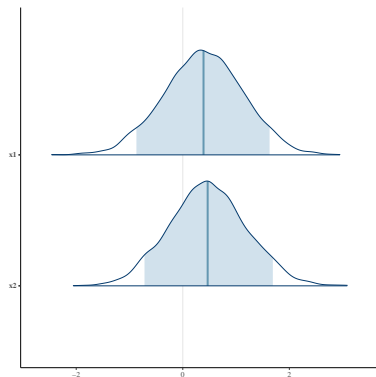
Marginal posterior intervals

Joint posterior density

rstanarm + bayesplot

see also Collinear demo

What if one is not clearly better than others?

What if one is not clearly better than others?

- Continuous expansion including all models?
  - and then analyse the posterior distribution directly
    avehtari.github.io/modelselection/betablockers.html
  - sparse priors like regularized horseshoe prior instead
    of variable selection
    video, refs and demos at avehtari.github.io/modelselection/

What if one is not clearly better than others?

- Continuous expansion including all models?
  - and then analyse the posterior distribution directly
    avehtari.github.io/modelselection/betablockers.html
  - sparse priors like regularized horseshoe prior instead
    of variable selection
    video, refs and demos at avehtari.github.io/modelselection/

- Model averaging with BMA or Bayesian stacking?
  mc-stan.org/loo/articles/loo2-example.html

What if one is not clearly better than others?

- Continuous expansion including all models?
    - and then analyse the posterior distribution directly
      avehtari.github.io/modelselection/betablockers.html
    - sparse priors like regularized horseshoe prior instead
      of variable selection
      video, refs and demos at avehtari.github.io/modelselection/

- Model averaging with BMA or Bayesian stacking?
  mc-stan.org/loo/articles/loo2-example.html

- In a nested case choose simpler if assuming some cost for
  extra parts?
  andrewgelman.com/2018/07/26/
  parsimonious-principle-vs-integration-uncertainties/

What if one is not clearly better than others?

- Continuous expansion including all models?
  - and then analyse the posterior distribution directly
    avehtari.github.io/modelselection/betablockers.html
  - sparse priors like regularized horseshoe prior instead
    of variable selection
    video, refs and demos at avehtari.github.io/modelselection/

- Model averaging with BMA or Bayesian stacking?
  mc-stan.org/loo/articles/loo2-example.html

- In a nested case choose simpler if assuming some cost for
  extra parts?
  andrewgelman.com/2018/07/26/
  parsimonious-principle-vs-integration-uncertainties/

- In a nested case choose more complex if you want to take
  into account all the uncertainties.
  andrewgelman.com/2018/07/26/
  parsimonious-principle-vs-integration-uncertainties/

Model averaging

- Prefer continuous model expansion

Model averaging

- Prefer continuous model expansion

- If needed integrate over the model space = model averaging

Model averaging

- Prefer continuous model expansion
- If needed integrate over the model space = model averaging
- Bayesian stacking may work better than BMA in case of misspecified models or small data
  - Yao, Vehtari, Simpson, & Gelman (2018)

Cross-validation and model selection

- Cross-validation can be used for model selection if
  - small number of models
  - the difference between models is clear

Cross-validation and model selection

- Cross-validation can be used for model selection if
  - small number of models
  - the difference between models is clear

- Do not use cross-validation to choose from a large set of models
  - selection process can lead to severe overfitting

Cross-validation and model selection

- Cross-validation can be used for model selection if
  - small number of models
  - the difference between models is clear

- Do not use cross-validation to choose from a large set of models
  - selection process can lead to severe overfitting

- Overfitting in selection process is not unique for cross-validation

Selection induced bias and overfitting

- Selection induced bias in cross-validation
  - same data is used to assess the performance and make the selection
  - the selected model fits more to the data
  - the CV estimate for the selected model is biased
  - recognized already, e.g., by Stone (1974)

Selection induced bias and overfitting

- Selection induced bias in cross-validation
  - same data is used to assess the performance and make the selection
  - the selected model fits more to the data
  - the CV estimate for the selected model is biased
  - recognized already, e.g., by Stone (1974)

- Performance of the selection process itself can be assessed using two level cross-validation, but it does not help choosing better models

Selection induced bias and overfitting

- Selection induced bias in cross-validation
  - same data is used to assess the performance and make the selection
  - the selected model fits more to the data
  - the CV estimate for the selected model is biased
 - recognized already, e.g., by Stone (1974)

- Performance of the selection process itself can be assessed using two level cross-validation, but it does not help choosing better models

- Bigger problem if there is a large number of models as in covariate selection

Selection induced bias in variable selection

Selection induced bias in variable selection

Take-home messages

- It's good to think predictions of observables, because observables are the only ones we can observe

- Cross-validation can simulate predicting and observing new data

- Cross-validation is good if you don't trust your model

- Different variants of cross-validation are useful in different scenarios

- Cross-validation has high variance, and **if** you trust your model you can beat cross-validation in accuracy

Take-home messages

- It's good to think predictions of observables, because observables are the only ones we can observe

- Cross-validation can simulate predicting and observing new data

- Cross-validation is good if you don't trust your model

- Different variants of cross-validation are useful in different scenarios

- Cross-validation has high variance, and **if** you trust your model you can beat cross-validation in accuracy

Take-home messages

- It's good to think predictions of observables, because observables are the only ones we can observe

- Cross-validation can simulate predicting and observing new data

- Cross-validation is good if you don't trust your model

- Different variants of cross-validation are useful in different scenarios

- Cross-validation has high variance, and **if** you trust your model you can beat cross-validation in accuracy

Take-home messages

- It's good to think predictions of observables, because observables are the only ones we can observe

- Cross-validation can simulate predicting and observing new data

- Cross-validation is good if you don't trust your model

- Different variants of cross-validation are useful in different scenarios

- Cross-validation has high variance, and **if** you trust your model you can beat cross-validation in accuracy

Take-home messages

- It's good to think predictions of observables, because observables are the only ones we can observe

- Cross-validation can simulate predicting and observing new data

- Cross-validation is good if you don't trust your model

- Different variants of cross-validation are useful in different scenarios

- Cross-validation has high variance, and **if** you trust your model you can beat cross-validation in accuracy